

Device driver with non-blocking read/write



Problem Statement:

We are writing a driver `/dev/iitpipe` with non-blocking read/write. Input of `iitpipe 0` is connected to output of `iitpipe1` and vice versa. So that the read write between the two pipelines goes into a buffer space .

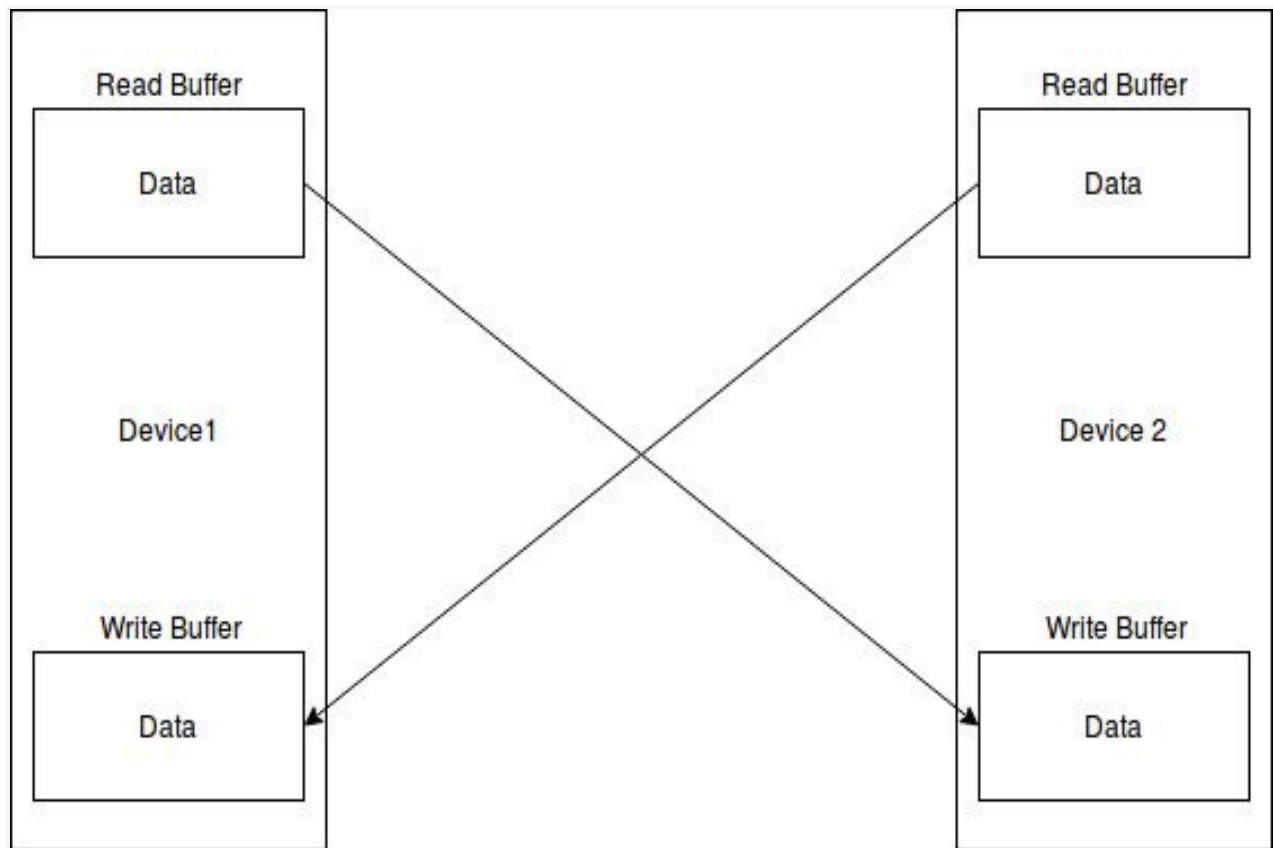
On `write()` the data is written into the pipe after a delay. We use a kernel timer or workqueue. The write delay should be controllable get and set via `ioctl()`. We use dynamic memory allocation.

Goals of our project:

Measuring the throughput of our driver with read/write buffer size set to 1 B, 1 KB, 1 MB and with delay set to 0, 1 ms, 1 sec.

****** Assuming the total data transferred be at least 100 MB in the 0 delay case and lower in other cases.

System design:



We have two device files with read/write, open/close which act as devices. The Read and Write buffers associated with the devices have the queue data structure.

```
struct circ_queue_node {  
    char byte;  
    long timestamp;  
};
```

Queue API:

circ_queue circ_queue_init(size_t size): Initializes the queue to given size

circ_queue_count(struct circ_queue queue): Counts the length of the queue (i.e nodes)

circ_queue_push(struct circ_queue *queue, struct circ_queue_node node): pushes the node into the queue.

circ_queue_peek(struct circ_queue *queue, struct circ_queue_node *node): outputs the first node of the queue.

circ_queue_pop(struct circ_queue *queue, struct circ_queue_node *node): pops and outputs the node from the queue.

circ_queue_free(struct circ_queue queue): free's the memory allocated to the queue .

Initialization:

We implemented the device driver as a module. When the module loads the it calls a `__init iitpipe_init(void)` using `module_init(iitpipe_init)` in which we first register the major number to the devices and create a class for the devices and then initialize the read/write buffers of the devices using `circ_queue circ_queue_init(size_t size)` function where size as default size().

How the data is read into Read Buffer ?

When we get the stream of data we store it first in a string using inbuilt functions `copy_from_user` which copies the string which comes as a argument into a local `char*` memory and then iteratively we push character by character into the read queue of the device.

And we free the local `char*` memory we created.

Timer-Function:

We have a timer function which calls the `update_queue` function which checks the time-stamp of the top element of the read buffer of a device and based on delay condition if satisfied it starts writing into the write buffer of the other device.

How IOCTL works?

In our project the `ioctl` is implemented as `my_ioctl` using `switch` which itself is a device file.

Based on the cases it does three tasks:

DELAY_GET : get information of current queues size and their delays parameters.

DELAY_SET : re-initialize the current queues dropping all the data it had initially and setting up the new parameters.

DELAY_CLR : All the memory allocated for the queues is cleared.

Test Cases/Performance results:

Buffer Size: 1 B

Delay 0ms - 110 bytes/s

Delay 1ms - same as above

Delay 1s - oscillates b/w 0 and 1

Buffer Size: 1 KB

Delay 0ms - 230 KB/s and oscillates dips down to 100KB/s some times

Delay 1ms - same as above

Delay 1s - oscillates b/w 1 and 0KB/s

Buffer Size: 1 MB

Delay 0ms - 17.5 MB/s 24MB/s

Delay 1ms - 17.5 MB/s 24MB/s

Delay 1s - Highly erratic and oscillates b/w 1 and 0 MB/s

Conclusion:

Learnt how device drivers work at kernel level. We were targeted to write at 100MB/s and we've achieved only 20MB/s.